

Seminar Series Neural Networks for Finance

Lecture 1

Aim

From linear to non-linear theory: towards an understanding of supervised learning

15-02-2024

Balint Negyesi (b.negyesi@tudelft.nl)

with the help of Kristoffer Andersson (kristoffer.andersson@cwi.nl)

Agenda

- ① General introduction to machine learning
- ② From statistics to machine learning: classical statistical inference
- ③ From linear to non-linear theory: regression
 - OLS regression
 - Non-linear regression
 - Generalization properties
- ④ From linear to non-linear theory: classification
 - Motivating example: Iris data set
 - Logistic regression
 - Other classification approaches
- ⑤ Non-linear theory: ANNs and DNNs
 - Single Layer Perceptron
 - Building blocks of neural networks
 - Activations
 - DNN
- ⑥ Teaser: capacity and optimization

General introduction to machine learning

Machine learning - General

Machine learning is a collection of techniques to "learn" a relationship between an *input space* A and an *output space* B . Often highly automated and with a statistical approach.

Example 1: (binary classification)

We want to classify images of cats and dogs. Assuming equal format of each image A can be a space of $X \times Y$ —matrices representing pixels of images and $B = \{\text{"cat"}, \text{"dog"}\}$.

Example 2: (function approximation)

We want to approximate $\mathbb{R} \ni x \mapsto f(x) \in [0, 1]$. Then, clearly $A = \mathbb{R}$ and $B = [0, 1]$.

Machine learning - General

Denote by H the true relation we want to learn, for $x \in A$, $H(x) \in B$ is the "truth". The aim for a machine learning algorithm is to achieve an accurate representation $\Phi \approx H$ in some sense. What do we need?

- A measure to define accuracy of our approximation $L(\Phi) = d(H, \Phi)$, referred to as the *loss function* (d refers to a non-specified metric).
- A collection of samples, or *features* $\mathbf{x}_M = (x_1, x_2, \dots, x_M)$ from A , i.e., $x_i \in A$.
- An empirical error measure, denoted $\bar{L}(\Phi|\mathbf{x}_M)$, satisfying $\lim_{M \rightarrow \infty} \bar{L}(\Phi|\mathbf{x}_M) = L(\Phi)$. Typically referred to as the *empirical loss function*.
- An optimization algorithm to approximately solve $\inf_{\Phi \in \mathcal{H}} \bar{L}(\Phi)$ in a pre-specified space \mathcal{H} . Typically referred to as *training*.

Machine learning - General scheme

Schematic scheme (some repetition of previous slide)

Given data: \mathbf{x}_M of size M (could also have labels in supervised learning, more on this later). Evaluation data $\bar{\mathbf{x}}_M$ of size \bar{M} .

① Training:

Find an approximate minimizer of the loss function Φ^* , *i.e.*, $\hat{\Phi} \approx \operatorname{argmin}_{\Phi \in \mathcal{H}} \bar{L}(\Phi|\mathbf{x}_M)$. Training or optimization is performed in a different way for each machine learning algorithm.

② Evaluation:

This step is to verify that the algorithm is able to generalize, *i.e.*, perform well beyond the training data. This is done by comparing $\bar{L}(\Phi|\mathbf{x}_M)$ and $\bar{L}(\Phi|\bar{\mathbf{x}}_M)$. If $\bar{L}(\Phi|\mathbf{x}_M) \ll \bar{L}(\Phi|\bar{\mathbf{x}}_M)$, then the problem might be *overfitting*, *i.e.*, the algorithm is fitted well against the training data but performs poorly on unseen data.

Machine learning - Different types of problems

Two main categories:

Supervised learning:

Each input element $x_i \in A$ comes with a corresponding output element $y_i \in B$. Implies that $H(x_i) = y_i$.

Unsupervised learning:

Only input elements are available during training. Need to rely on specific structure of the problem at hand to construct a loss function.

Supervised learning:

- Training data comes in pairs (x_i, y_i) such that $H(x_i) = y_i$, for $i = 1, 2, \dots, M$, where y_1, y_2, \dots, y_M are referred to as *labels*.
- As empirical loss function $L(\Phi|\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M d(\Phi(x_m), y_m)$ can be used.

Example: (Binary classification - American options)

At a certain point in time, given the price of the underlying, should the option be held or exercised? Optimal stopping policy.

Machine learning - Supervised learning Example

Input data: x_1, \dots, x_M representing M states of the underlying asset.

Labels: y_1, \dots, y_M taking on values in $\{0, 1\}$ where "0" means sub-optimal to exercise option and "1" means it is optimal to exercise option.

The aim is to approximate $H : \mathbb{R}_+ \rightarrow \{0, 1\}$ with Φ .

Loss function: $L(\Phi|\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M d(\Phi(x_m), y_m)$, where $d(a, b)$ could be $|a - b|$ but usually more sophisticated to have a smoother loss function.

Unsupervised learning:

- Training data consists of only input data x_1, \dots, x_M and the aim is to approximate H by Φ without being able to observe $H(x_i)$.
- We need to come up with an empirical loss function
- As empirical loss function $L(\Phi|\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M F(\Phi(x_m))$, where F is a certain *functional* such that if $L(\Phi|\mathbf{x})$ is minimized (or maximized), then Φ is a good approximator of H (at least for states around x_1, \dots, x_M).

Example: (Option valuation - European options)

At a certain point in time, given the price of the underlying, what is the value of a European option? Solving the pricing PDE.

Machine learning - Unsupervised learning Example

Pricing PDE:

$$\begin{cases} \frac{\partial V}{\partial t} = f(t, x, V, \frac{\partial V}{\partial x}, \frac{\partial^2 V}{\partial x^2}) & (t, x) \in [0, T] \times \Omega = (x_{\min}, x_{\max}), \\ V(t, x) = f(t, x), & (t, x) \in [0, T] \times \bar{\Omega}, \\ V(T, x) = g(x), & x \in \Omega = \{x_{\min}, x_{\max}\}. \end{cases}$$

Here f is determined by the dynamics of the underlying asset, g is the pay-off function and the boundary conditions are set up to be "asymptotically reasonable".

For instance, the Black–Scholes model gives f linear in $V, \frac{\partial V}{\partial x}, \frac{\partial^2 V}{\partial x^2}$ making the PDE linear. Other models may yield a semi-linear or even a fully non-linear PDE.

Machine learning - Unsupervised learning Example

We do not know the PDE solution apriori and therefore, have no access to labels as in supervised learning. Need to create a loss function which uses the structure from the PDE.

$$L(\Phi) = \text{MSE}_{\text{inner}} + \text{MSE}_{\text{TC}} + \text{MSE}_{\text{BC}},$$

where

$$\begin{cases} \text{MSE}_{\text{inner}} = \frac{1}{M_{\text{inner}}} \sum_{m=1}^{M_{\text{inner}}} \left| \frac{\partial \Phi(t_m, x_m)}{\partial t} - f(t_m, x_m, \Phi, \frac{\partial \Phi}{\partial x}, \frac{\partial^2 \Phi}{\partial x^2}) \right|^2 \\ \text{MSE}_{\text{TC}} = \frac{1}{M_{\text{TC}}} \sum_{m=1}^{M_{\text{IC}}} |g(x_m) - \Phi(T, x_m)|^2, \\ \text{MSE}_{\text{BC}} = \frac{1}{M_{\text{BC}}} \sum_{m=1}^{M_{\text{BC}}} |f(t_m, x_m) - \Phi(t_m, x_m)|^2. \end{cases}$$

Machine learning – Other categories of tasks

- Anomaly detection
- Synthesis and sampling
- Denoising
- Density estimations
- ...

From statistics to machine learning: classical statistical inference

From statistics to machine learning: classical statistical inference

"Machine learning is essentially a form of applied statistics with increased emphasis on the use of computers to statistically estimate complicated functions and a decreased emphasis on proving confidence intervals around these functions[...]"

Goodfellow et al.: Deep Learning (2016), MIT Press

Statistical Estimates

- Statistics is about extracting as much information as possible from measurement data, e.g. *fitting functions, estimating finite/infinite parameters, distributions, assessing confidence intervals around estimates*
- In the most general, measure theoretic statement: a statistical model is a family of probability measures $\{\mathbb{P}_\theta, \theta \in \Omega\}$ on a measurable space $(\mathcal{X}, \mathcal{B})$ acting on the distribution of data X
 - Ω **finite dimensional**: **parametric models**
 - Ω infinite dimensional: non-parametric models
- In this course we will focus on **parametric models**

Statistical Point Estimates

- Two main schools of statistical inference:
 - ① Classical/Frequentist: the true value of a parameter θ is fixed but not known; the *point estimate* $\hat{\theta}$ is a random variable as a function of the randomly generated data
 - ② Bayesian: data is directly observed thus cannot be random; the true value of the parameter θ is unknown therefore uncertain and thus can be modeled as a random variable
- Frequentist: summarizes data into a single value of the parameter set $\hat{\theta}$ and then makes predictions
- Bayesian: encompasses prior knowledge into a **prior** probability distribution $p(\theta)$ which is then used to make predictions through the Bayes theorem $p(\theta|x) = \frac{p(x|\theta)p(\theta)}{p(x)}$
- Through most of the course we will take the frequentists' perspective, however when studying **General Adversarial Networks** (GANs) we will rely on the **Bayesian approach**

Estimators, bias, variance

- Given data $\{x_1, \dots, x_m\}$, a point estimator is a function which maps the data to an element of the parameter space

$$\hat{\theta}_m := g(x_1, \dots, x_m)$$

- Example: $x_i \in \mathbb{R}$, the statistical model consists of fitting a line on the data points
- Bias:** $\text{bias}(\hat{\theta}) := \mathbb{E} [\hat{\theta}] - \theta$, where the expectation is taken over the data-generating distribution, $\text{bias}(\hat{\theta}_m) = 0 \implies$ **unbiased** estimator
- Variance:** $\text{Var} [\hat{\theta}_m]$
- The **mean-squared error (MSE)** defines a trade-off between the minimization of the two

$$\text{MSE}(\hat{\theta}_m) = \mathbb{E} [(\hat{\theta}_m - \theta)^2] = \text{bias}^2(\hat{\theta}_m) + \text{Var} [\hat{\theta}_m]$$

Notions of generalization

$$\text{MSE}(\hat{\theta}_m) = \mathbb{E} \left[(\hat{\theta}_m - \theta)^2 \right] = \text{bias}^2(\hat{\theta}_m) + \text{Var} \left[\hat{\theta}_m \right]$$

- Notions of generalization: capacity, underfitting, overfitting
 - 1 underfitting: small capacity leads to low variance, large bias
 - 2 overfitting: increased capacity, decreases bias but increases variance

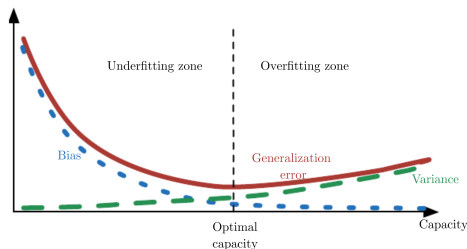


Figure: Goodfellow et al.: Deep Learning (2016), MIT Press [Figure 5.6, pg. 128]

Maximum Likelihood Principle

- **Maximum Likelihood Estimator (MLE)** principle: given a data generating distribution $\mathbb{P}_{\text{data}}(x)$ and a statistical model $\mathbb{P}_{\text{model}}(x|\theta)$
$$\hat{\theta}_{\text{ML}} := \arg \max_{\theta} \mathbb{P}_{\text{model}}(X|\theta) \quad (1)$$
- Under the assumption of **independent, identically distributed (iid)** data samples $X = \{x_1, \dots, x_M\}$, this can be written $\hat{\theta}_{\text{ML}} := \arg \max_{\theta} \prod_{m=1}^M \mathbb{P}_{\text{model}}(x_m|\theta)$

From linear to non-linear theory: regression

Ordinary Least-Squares Regression

- We are given a data set $\{x_{m1}, \dots, x_{md}; y_m\}_{m=1}^M$ of M measurements drawn from the same distribution
- The elements of the vector $\mathbf{x}_m := (x_{m1}, \dots, x_{md}) \in \mathbb{R}^d$ are called the *explanatory/independent* variables; $y_m \in \mathbb{R}$ is the response/dependent variable
- We assume a *linear relation with noise* between the *fixed* measurements of independent variables and the response variable

$$y = \mathbf{x}^T \boldsymbol{\beta} + \epsilon \quad (2)$$

with some **parameter** vector $\boldsymbol{\beta} \in \mathbb{R}^d$

- the *noise term* ϵ incorporates the residual influence on y not captured by the regressors (measurement errors in y , missing regressors, etc.)

Ordinary Least-Squares Regression

Goal

The goal of Ordinary Least-Squares (OLS) regression is to find a(/the?) set of regression coefficients $\beta = (\beta_1, \dots, \beta_d)$ such that the Euclidean norm of the noise term is minimized

$$\hat{\beta} := \arg \min_{\beta} \|\epsilon\|_2^2 = \arg \min_{\beta} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 = \sum_{m=1}^M \left| y_m - \mathbf{x}_m^T \beta \right|^2, \quad (3)$$

where $\mathbf{X} = \begin{pmatrix} x_{11} & \dots & x_{1d} \\ \dots & \dots & \dots \\ x_{m1} & \dots & x_{md} \end{pmatrix}$, $\mathbf{y} = (y_1, \dots, y_m)$, $\epsilon = (\epsilon_1, \dots, \epsilon_m)$

Ordinary Least-Squares Regression

- Denote mean-squared *loss*, subject to the measurements $\mathbf{D} := (\mathbf{X}, \mathbf{y})$ by $L(\beta|\mathbf{D}) = \|\mathbf{y} - \mathbf{X}\beta\|_2^2$ which is equal to

$$L(\beta|\mathbf{D}) = \mathbf{y}\mathbf{y}^T - \mathbf{y}^T\mathbf{X}\beta - \beta^T\mathbf{X}^T\mathbf{y} + \beta^T\mathbf{X}^T\mathbf{X}\beta$$

- Due to the **convexity** of $\beta \mapsto L(\beta|\mathbf{D})$ the unique global minimum is found where the gradient vanishes

$$\mathbf{0} = \nabla_{\beta} L(\hat{\beta}|\mathbf{D}) = -2\mathbf{y}^T\mathbf{X} + 2\hat{\beta}^T\mathbf{X}^T\mathbf{X}$$

Unique Solution

The unique minimizer of (3) is given by

$$\hat{\beta} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}, \quad (4)$$

and the corresponding approximations

$$\hat{\mathbf{y}} := \mathbf{X}\hat{\beta} \quad (5)$$

Key Remarks

- Under standard assumptions (linearity, independence of ϵ_m 's, homoscedasticity – the variance of ϵ_m does not depend on the values \mathbf{x}_m , lack of *perfect multicollinearity* – \mathbf{X} has full column rank d) the OLS problem admits to a **unique, closed-form solution**
- The intercept (constant) can be incorporated into the model as a zero'th dimension
- The distribution of \mathbf{X} has a great impact on the accuracy of $\hat{\beta}$ – assumes high-quality sampling methods
- The OLS estimate coincides with the Maximum Likelihood Estimate (MLE)

Non-linear regression

- Same setting
 - data coming from a fixed *data generating* distribution: $\{x_{m1}, \dots, x_{md}; y_m\}_{m=1}^M$ of M measurements drawn from the same distribution
 - explanatory and response variables $\mathbf{x}_m := (x_{m1}, \dots, x_{md}) \in \mathbb{R}^d$, $y_m \in \mathbb{R}$
- However, the assumptions of linearity established by (2) **no longer holds**
- Instead, we consider the following **non-linear extension** of the model

$$y = H(\mathbf{x}) + \epsilon \quad (6)$$

where $H : \mathbb{R}^d \rightarrow \mathbb{R}$ is a deterministic mapping

Question?

What could be an appropriate family of statistical models $\{\phi(x|\theta) : \mathbb{R}^d \rightarrow \mathbb{R} : \theta \in \mathbb{R}^p\}$ which spans a "large enough" space for all potential H s?

- In what follows we will study classes of statistical models which are suitable for specific tasks
- For now, let us briefly assume we have a family of "**black box**" statistical models of the form $\{\phi(x|\theta) : p \in \mathbb{N}, \theta \in \mathbb{R}^p\}$ which possess this quality
- Within such a family of statistical models, the question still remains: *how do we choose the "right" number of p parameters?*

Generalization Properties

- Here p is called a **hyperparameter** of the model which needs to be chosen carefully
- The ultimate goal of the regression problem is not to fit the *training data* $\{x_{m1}, \dots, x_{md}; y_m\}_{m=1}^M$ with arbitrary accuracy but rather to **approximate the relation** (6)
- In fact, we seek to find a set of parameters $\{\hat{\theta} \in \mathbb{R}^p : p \in \mathbb{N}\}$ such that $\phi(x|\hat{\theta}) \approx H(x)$ for any x **drawn from the data generating distribution** – regardless whether it is in the training data or not
- trade-off: representational capacity naturally grows with p , however generalization capability decreases – recall Figure 1

Hyperparameter Selection: Under- and Overfitting

In some applications one has direct access to the data generating distribution (e.g.: Monte Carlo option pricing), in most cases however we are only provided with a finite sample

$$\mathbf{D} = \{(\mathbf{x}_m, y_m)\}_{m=1}^M$$

Question

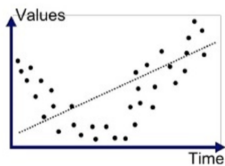
Question: in case of the latter, how does one ensure *unbiased* estimates, with good generalization properties?

- Idea: split the data set \mathbf{D} into a partition, a subset for training $\mathbf{D}_{train} \subset \mathbf{D}$ and one for testing $\mathbf{D}_{test} \subset \mathbf{D}$, such that $\mathbf{D}_{train} \cap \mathbf{D}_{test} = \emptyset$ and $\mathbf{D}_{train} \cup \mathbf{D}_{test} = \mathbf{D}$.
- We choose multiple p 's and *optimize* their corresponding statistical models with a **black box** optimization method (soon explained) according to an appropriately selected *loss function* (e.g. MSE)

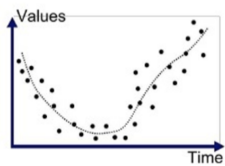
Hyperparameter Selection: Under- and Overfitting

- Out of the optimal parameter sets, the one with lowest loss generalizes best. However, the error measure remains inaccurate to future predictions due to overfitting on the test set.
- Partitioning into three subsets: training data $\mathbf{D}_{train} \subset \mathbf{D}$, validation data $\mathbf{D}_{val} \subset \mathbf{D}$ and test data $\mathbf{D}_{test} \subset \mathbf{D}$, such that $\mathbf{D}_{train} \cup \mathbf{D}_{val} \cup \mathbf{D}_{test} = \mathbf{D}$ and $\mathbf{D}_{train} \cap \mathbf{D}_{val} \cap \mathbf{D}_{test} = \emptyset$ resolves this problem
- First, multiple parameter sets are trained on the training set.
- Second, the best performing set of parameters is chosen from the validation set.
- At last, the error measure is calculated on the test set.
- This method gives an unbiased estimate of the error measure on future data.
- Similarly to **cross validation**

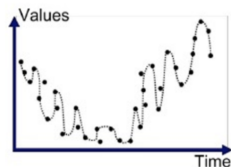
Summary: Under- and Overfitting



Underfitted



Good Fit/Robust



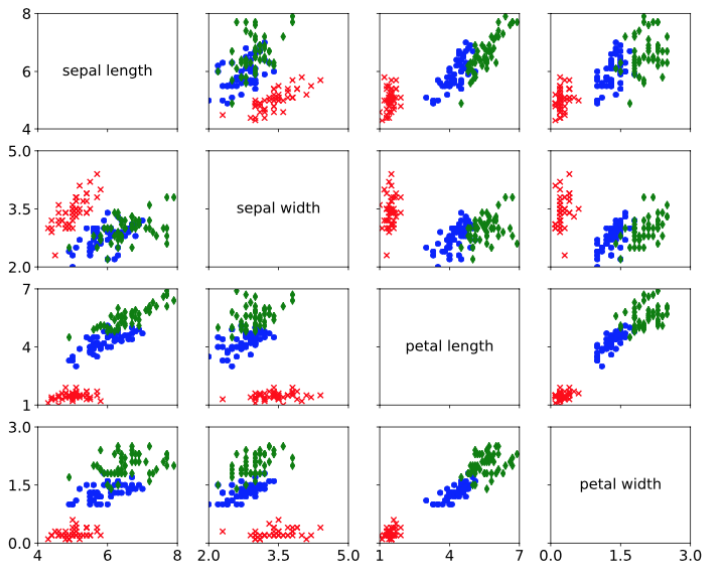
Overfitted

From linear to non-linear theory: classification

Motivating Example: Iris data set

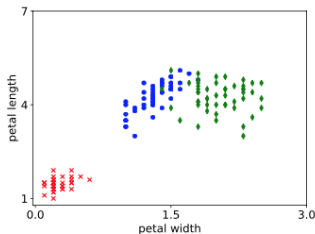
- The Iris data set is a multivariate set with 50 entries for each species. There are three species, consisting of four features: sepal length, sepal width, petal length and petal width.
- The four features can be displayed in a scatter plot
- Determining categories by using their features is called **classification**
- Setosa is blue, Versicolor is red, and Virginica is green colored

Motivating Example: Iris data set

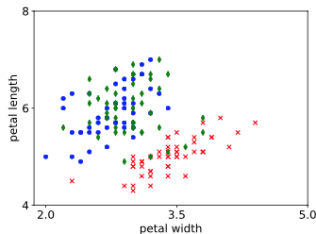


Motivating Example: Iris data set

- See, left figure below, the Setosa species can be distinguished from the others, by using the petal length.



(a) Zoom 1



(b) Zoom 2

- In mathematical terms, it can be stated that the Setosa is **linearly separable**.
- Two sets $X_1 \subset X$ and $X_2 \subset X$ are called **linearly separable** if their convex hulls are disjoint.

Motivating Example: Iris data set

- Consider the set X_1 , the petal lengths of the Versicolor and the Virginica species and the set X_2 , the petal lengths of the Setosa species. In this case

$$\text{convhull}(X_1) = \{x | x \in [\min(X_1), \max(X_1)]\} = [3.0, 6.9]$$

and

$$\text{convhull}(X_2) = \{x | x \in [\min(X_2), \max(X_2)]\} = [1.0, 1.9]$$

- Clearly, $\text{convhull}(X_1) \cap \text{convhull}(X_2) = \emptyset$, therefore X_1 and X_2 are linearly separable, and thus it is possible to classify the Setosa species in the Iris data set

Motivating Example: Iris data set

- From the scatter plot at the right side of the last figure, the species are not easily distinguishable
- By creating convex hulls of the individual features, the two species are not linearly separable for the individual features
- It is impossible to classify the species using linear classifiers deterministically. We could create (hyper-)planes such that a minimal number of data entries is incorrectly separated
- This gives a statistically optimal result for the data set on which the classifier is trained
- **Support Vector Machines (SVM)** designed to find such optimal hyperplanes

Standard models of classification

Standard Models of Classification

- We are given a **binary** response variable $Y \in \{0, 1\}$; and a vector of real-valued regressors \mathbf{x} which are assumed to influence Y – *multinomial extensions are straightforward*
- We obtained a data set of measurements $\{\mathbf{x} = (x_{m1}, \dots, x_{md}), y_m\}_{m=1}^M$
- Formally we describe the outcome of each measurement with a Bernoulli distribution with *unobserved* probability p_m

$$Y_m | x_{m1}, \dots, x_{md} \sim \text{Bernoulli}(p_m)$$

- Define the linear *latent variable model* with an *auxiliary random variable* \tilde{Y} such that

$$\tilde{Y} = \mathbf{x}^T \boldsymbol{\beta} + \epsilon, \quad (7)$$

where ϵ follows a given **symmetric, centered** distribution

Standard Models of Classification

- Thereby Y can be interpreted as the indicator function

$$Y = \begin{cases} 1, & \tilde{Y} > 0 \\ 0, & \tilde{Y} \leq 0 \end{cases}$$

- By the symmetricity of the distribution of the noise term we have

$$\mathbb{P}[Y = 1|\mathbf{x}] = \mathbb{P}[\tilde{Y} > 0|\mathbf{x}] = \mathbb{P}[\epsilon > -\mathbf{x}^T\boldsymbol{\beta}] = \text{CDF}(\mathbf{x}^T\boldsymbol{\beta})$$

- The corresponding estimates are then defined by setting the distribution of the noise term, two classical choices
 - ➊ **Probit regression:** $\epsilon \sim \mathcal{N}(0,1)$ which implies

$$\mathbb{P}[Y = 1|\mathbf{x}] = \Phi(\mathbf{x}^T\boldsymbol{\beta})$$

- ➋ **Logistic regression:** ϵ follows a *standard logistic distribution* whose CDF is given by $\frac{1}{1+e^{-x}}$

$$\mathbb{P}[Y = 1|\mathbf{x}] = \frac{1}{1 + e^{-\mathbf{x}^T\boldsymbol{\beta}}}$$

Logistic Regression

- Logistic regression is more common in machine learning practices due to the convenience of a closed form CDF
- So far we have a prediction for the probability of the m 'th measurement taking the value 1

$$\mathbb{P}[Y_m = 1 | \mathbf{x}_m, \beta] = p(\mathbf{x}_m | \beta) = \frac{1}{1 + e^{-\mathbf{x}_m^T \beta}}$$

- Notice that the prediction of each class is no longer a linear mapping of X ; only in case of the latent variable
- In fact, the model is not linear in the predicted probabilities but rather in the so called log-**odds**'es

$$\text{odds}(\mathbf{x}) := \frac{p(\mathbf{x} | \beta)}{1 - p(\mathbf{x} | \beta)} = e^{-\mathbf{x}^T \beta}$$

- This establishes the direct relation between standard linear and logistic regressions

Logistic Regression

- Due to the Bernoulli condition, under the assumption of independent samples the corresponding likelihood function is of the form

$$\mathcal{L}(\beta|\mathbf{X}, \mathbf{y}) = \prod_{m=1}^M \mathbb{P}[y_m|\mathbf{x}_m, \beta] = \prod_{m=1}^M p(\mathbf{x}_m|\beta)^{y_m} (1-p(\mathbf{x}_m|\beta))^{1-y_m}$$

- Unlike in OLS regression the model **does not have a closed-form** (maximum likelihood) solution
- To formulate an equivalent loss minimization problem we define $L(\beta|\mathbf{X}, \mathbf{y}) := -\log \mathcal{L}(\beta|\mathbf{X}, \mathbf{y})$ leading to the **cross-entropy loss**

$$L(\beta|\mathbf{X}, \mathbf{y}) = - \sum_{m=1}^M y_m \log(p(\mathbf{x}_m|\beta)) + (1 - y_m)(1 - p(\mathbf{x}_m|\beta)) \quad (8)$$

$$L(\beta|\mathbf{X}, \mathbf{y}) = - \sum_{m=1}^M y_m \log(p(\mathbf{x}_m|\beta)) + (1 - y_m)(1 - p(\mathbf{x}_m|\beta)) \quad (9)$$

- There is **no closed-form** solution for the minimizer
- Iterative root searching methods can be applied to gather an estimate for the minimizer

$$\hat{\beta} := \arg \min_{\beta} L(\beta|\mathbf{X}, \mathbf{y})$$

- For instance: Newton's method, *(stochastic) gradient descent coming soon...*
- The gradient of the cross-entropy loss is equal to that of MSE

$$\nabla_{\beta} L(\beta|\mathbf{X}, \mathbf{y}) = \mathbf{X}(\hat{\mathbf{y}} - \mathbf{y})$$

Other classification approaches

- clustering
- k -nearest neighbours
- isolation forests
- decision trees
- support vector machines (SVM)
- naive Bayes classifier
- ...

Non-linear theory: Artificial Neural Networks and Deep Neural Networks

From linear to non-linear mappings

- Linear models have the obvious defect of model capacity only spanning the space of linear mappings
- We wish to extend and represent non-linear functions of some vector input $\mathbf{x} \mapsto H(\mathbf{x})$
- Idea: apply a linear model to not \mathbf{x} itself but rather to a non-linear, vector-valued transformation of it
 $\phi(\mathbf{x}|\theta) : \mathbb{R}^d \rightarrow \mathbb{R}^p$
- The resulting approximations of $H(\mathbf{x})$ read as

$$\Phi(\mathbf{x}|\Theta := (\theta, \mathbf{w})) = \varphi^T(\mathbf{x}|\theta)\mathbf{w},$$

with $\mathbf{w} \in \mathbb{R}^p$

Question

How do we choose the family of $\Phi(\mathbf{x}|\Theta)$?

The Single Layer Perceptron Model

- **Single Layer Perceptron (SLP)** is the simplest neural network model originally designed to solve the binary classification problem

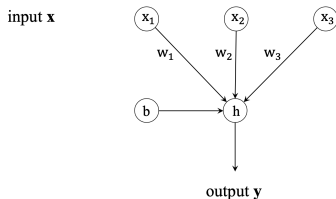


Figure: Illustration of an SLP

- Inputs (written as \mathbf{x}) are passed through a weighted, directed graph.
- A function h corresponds to the non-linear mapping $\phi(\mathbf{x}|\theta)$ from above

The Single Layer Perceptron Model

- SLP has $d + 1$ parameters: a d -dimensional vector of *weights* $\mathbf{w} \in \mathbb{R}^d$ connecting the node of h to each input feature; and a *bias* denoted by $b \in \mathbb{R}$
- The non-linearity $h_b : \mathbb{R} \rightarrow \{-1, 1\}$ is defined by the following step function with threshold b

$$h_b(z) := \begin{cases} 1, & z > b, \\ -1, & z \leq b, \end{cases} \quad (10)$$

which can be interpreted as **true** or **false**

- The input of the non-linearity is a scalar defined by the inner product of the input vector $\mathbf{x} \in \mathbb{R}^d$ and the weights: $\mathbf{w}^T \mathbf{x}$
- Thus the mapping of SLP is given by the following *sequence of compositions*

$$\text{SLP}(\mathbf{x} | \theta := (\mathbf{w}, b)) = \begin{cases} 1, & \mathbf{w}^T \mathbf{x} > b, \\ -1, & \mathbf{w}^T \mathbf{x} \leq b. \end{cases} \quad (11)$$

Training: Optimization Analytically

Question

How to find an **optimal** set of parameters properly assigning a binary classification label?

- For linearly separable data, an algorithm optimizes/trains the parameter vector $\theta = (-b, \mathbf{w})$ such that the output of the perceptron corresponds to the correct class for all data entries in X .
- The algorithm starts by initializing a random parameter vector and a convergence test variable. It then picks a random sample from the data set and tests whether the sign of the weight vector assigns the correct label to the random sample.
- If an incorrect label is assigned, it updates the weight such that the argument of the sign function moves to the correct label.

Training: Optimization Analytically

- For notational convenience we attach the bias node to the input as a 0'th (constant) dimension $\mathbf{x}' = (1, \mathbf{x}) \in \mathbb{R}^{d+1}$, and assign it $w_0 = -b$ weight. These thresholds are equivalent

$$\boldsymbol{\theta}^T \mathbf{x}' \geq 0 \iff \mathbf{w}^T \mathbf{x} \geq b \implies h_0(\boldsymbol{\theta}^T \mathbf{x}') \equiv h_b(\mathbf{w}^T \mathbf{x}).$$

- Consider the case of a positive label in y and a negative prediction label, the weight is updated by adding the data \mathbf{x}' . To see why this works, consider:

$$h_0(\boldsymbol{\theta}^T \mathbf{x}') = -1. \quad (12)$$

Add \mathbf{x}' to $\boldsymbol{\theta}$ and use the linearity of the inner product:

$$h_0((\boldsymbol{\theta} + \mathbf{x}')^T \mathbf{x}') = h_0(\boldsymbol{\theta}^T \mathbf{x}' + \mathbf{x}'^T \mathbf{x}') \quad (13)$$

- By definition, for the inner product we have $\mathbf{x}'^T \mathbf{x}' \geq 0$. Since $\mathbf{x}' \neq 0$ from (12), $\mathbf{x}'^T \mathbf{x} > 0$.
- Therefore the weight update improves the prediction

Perceptron Learning Algorithm for linearly separable data

- Let X be an $M \times (d + 1)$ matrix with M measurements, d features and a column for the bias (all values equal 1), such that two classes in the set are linearly separable; let length M vector $\mathbf{y} \in \{-1, 1\}^M$ indicate the class of a data entry of X .
- Choose a random $\boldsymbol{\theta} \in \mathbb{R}^{d+1}$. Set $conv = 1$
- while $conv$ do
 - Choose a random $r \in U(0, M)$
 - Set $\mathbf{x}' = X(r, \cdot)$
 - Set $y = \mathbf{y}(r)$
 - Compute $\hat{y} = \text{sign}(\boldsymbol{\theta}^T \mathbf{x}')$
- if $y = 1$ and $\hat{y} = -1$ then
 - Update $\boldsymbol{\theta} := \boldsymbol{\theta} + \mathbf{x}'$, endif
- if $y = -1$ and $\hat{y} = 1$ then
 - Update $\boldsymbol{\theta} := \boldsymbol{\theta} - \mathbf{x}'$, endif
- Convergence test: $conv := 0$ if $h_0(X^T \mathbf{w}) = \mathbf{y}$
- end while

Summary

- After each random sample, a convergence test is done by comparing all the signs of the samples in the data set with their corresponding labels.
- If not all data entries are classified correctly, the algorithm starts over by picking a new random sample, otherwise the algorithm has successfully found a set of weights which classifies the data, and terminates.
- To speed up the training time, one might consider using different algorithms (training of perceptrons can be translated to solving an interior point problem).

Logistic regression as an SLP

- In the previous SLP example we specified a step function to model non-linear behavior
- If one replaces $h : \mathbb{R} \rightarrow \mathbb{R}$ with the CDF of the standard logistic distribution

$$\text{CDF}(\boldsymbol{\theta}^T \mathbf{x}') = \frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{x}'}} ,$$

the resulting SLP scheme is equivalent to the logistic regression model

- The resulting equivalent SLP model can then be optimized in an identical fashion as in case of logistic regression: minimizing the cross-entropy loss with an iterative root searching algorithm

Multinomial Logistic regression

- One can easily extend the binary logistic regression to multinomial models of K outcomes $\{0, \dots, K-1\}$
- This is done by running $K-1$ many binary classifications with a fixed outcome as pivot

$$\frac{\mathbb{P}[Y_m = 0]}{\mathbb{P}[Y_m = K-1]} = e^{-\beta_1^T \mathbf{x}_m}, \quad \dots, \quad \frac{\mathbb{P}[Y_m = K-2]}{\mathbb{P}[Y_m = K-1]} = e^{-\beta_{K-1}^T \mathbf{x}_m},$$

where $\beta_k \in \mathbb{R}^d$

- The probabilities assigned to each class must sum up to 1:

$$\begin{aligned} \mathbb{P}[Y_m = K-1] &= 1 - \sum_{k=0}^{K-2} \mathbb{P}[Y_m = k] \\ \implies \mathbb{P}[Y_m = K-1] &= \frac{1}{1 + \sum_{k=0}^{K-2} e^{-\beta_k^T \mathbf{x}_m}}, \end{aligned}$$

where $\beta_k \in \mathbb{R}^{d+1}$ is the coefficients corresponding to the k 'th class

Multinomial Logistic regression

- The resulting multinomial logistic regression scheme assigns the following probabilities to class $k \leq K - 2$

$$\mathbb{P}[Y_m = k | \boldsymbol{\theta}, \mathbf{x}_m] = \frac{e^{-\boldsymbol{\beta}_k^T \mathbf{x}_m}}{1 + \sum_{k=0}^{K-2} e^{-\boldsymbol{\beta}_k^T \mathbf{x}_m}} \quad (14)$$

- The set of parameters $\boldsymbol{\theta} \in \mathbb{R}^{K-1 \times d+1}$ becomes a matrix of *weights* defining the mapping to each outcome class
- For the Iris classification problem, the vector consists of three elements, where vector $(1, 0, 0)$ is assigned to Setosa, $(0, 1, 0)$ is assigned to the class Versicolor and $(0, 0, 1)$ to Virginica

Example: XOR problem

- Example that shows why an SLP may be insufficient for simple classification and gives rise more refined non-linear models
- Consider the XOR-problem (exclusive OR relation)
- In this problem the data is not linearly separable. We have input $\mathbf{x} \in \mathbb{R}^2$ and a class $y \in \{-1, 1\}$ as shown below.

x_1	x_2	y
-1	-1	1
-1	1	-1
1	-1	-1
1	1	1

- **Impossible to classify this using a Single Layer Perceptron**

Building blocks of neural networks

Building blocks of neural networks

Composition of simple functions

- SLP is only capable of computing binary functions
- It nevertheless acts as a motivating *structure* for non-linear parametrizations
- Linear combinations of the input are passed to a non-linear function

Idea

Generalize this structure: build non-linear parametrizations as **hierarchical compositions** of non-linear and linear transformations → **neural networks**

Building blocks of composition: affine transformations

- Recall: in OLS $\hat{y} = \beta^T \mathbf{x}$ where $\beta \in \mathbb{R}^{d+1}$ is the parameter vector
- In the context of neural networks, affine mappings of the form

$$\mathbb{R}^d \ni \mathbf{x} \mapsto \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)} =: \mathbf{z}^{(1)} \in \mathbb{R}^{p_1} \quad (15)$$

are called the **linear layers**, which map their input \mathbf{x} to an affine combination with **weights** $\mathbf{W}^{(1)} \in \mathbb{R}^{p_1 \times d}$ and **biases** $\mathbf{b}^{(1)} \in \mathbb{R}^{p_1}$

- The total number of parameters is $\theta^{(1)} = (\mathbf{W}^{(1)}, \mathbf{b}^{(1)}) \in \mathbb{R}^{p_1 \times (d+1)}$
- The length of $\mathbf{z}^{(1)}$ is called the **width** of the layer
- The components of $\mathbf{z}^{(1)}$ are called **neurons**
- The number of neurons in a hidden layer is a *hyperparameter*, the wider it is set, the more parameters there are

Building blocks of composition: hidden layers

- The machine learning community often confuses the terms *linear* and *affine*. This is due to the fact that the constant bias (similar as in OLS or SLP) can be viewed as an additional constant dimension with weight vector \mathbf{b} .
- So far this is identical to the model of (vector-valued) OLS with *cast*: $\hat{\mathbf{y}} \leftarrow \mathbf{z}^{(1)}$ and $\beta \leftarrow \theta^{(1)}$
- Recall idea: apply a (non-)linear vector-valued mapping to the affine combinations of the input $\mathbf{w}^T \phi(\mathbf{z}^{(1)} | \theta^{(1)})$

$$\mathbb{R}^{p_1} \ni \mathbf{z}^{(1)} \mapsto \left(\varphi_1^{(1)}(z_1^{(1)}), \dots, \varphi_{p_1}^{(1)}(z_{p_1}^{(1)}) \right) =: \mathbf{a}^{(1)} \in \mathbb{R}^{p_1},$$

where $\varphi_i^{(1)} : \mathbb{R} \rightarrow \mathbb{R}$ are non-linear **activations**

- The composition of these transformations is called a **hidden layer** of a neural network

$$\mathbf{x} \mapsto \varphi_i^{(1)} \left(\sum_{j=1}^d W_{ij}^{(1)} x_j + b_i^{(1)} \right) =: a_i^{(1)}, \forall i = 1, \dots, p_1 \quad (16)$$

Building blocks of composition: output layer

- In above, p is a *hyperparameter* which, similarly to OLS, can be interpreted as the number of *basis functions* $\{x_1, \dots, x_d\}$ in a regression
- In order to approximate vector valued functions $H : \mathbb{R}^d \rightarrow \mathbb{R}^q$, one needs to transform the mapping of the hidden layer to the appropriate dimension of the *output*
- As before, we do this by an affine combination of the output with another set of *weights* and *biases*

$$\mathbb{R}^{p_1} \ni \mathbf{a}^{(1)} \mapsto \mathbf{W}^{(2)} \mathbf{a}^{(1)} + \mathbf{b}^{(2)} =: \mathbf{z}^{(2)} \in \mathbb{R}^q \quad (17)$$

with $\mathbf{W}^{(2)} \in \mathbb{R}^{q \times p_1}$, $\mathbf{b}^{(2)} \in \mathbb{R}^q$,
 $\boldsymbol{\theta}^{(2)} := (\mathbf{W}^{(2)}, \mathbf{b}^{(2)}) \in \mathbb{R}^{q \times (p_1 + 1)}$

- Finally, in order to support the most general (bounded, binary, ...) output, the **output layer** is a composition

$$\mathbb{R}^{p_1} \ni \mathbf{a}^{(1)} \mapsto \left(\varphi_1^{(2)}(z_1^{(2)}), \dots, \varphi_q^{(2)}(z_q^{(2)}) \right) =: \mathbf{a}^{(2)} \in \mathbb{R}^q$$

(Fully-Connected) Feedforward Artificial Neural Networks

- A **fully-connected, feedforward artificial neural network** is simply a **hierarchical composition** of the transformations above, mapping an input $\mathbf{x} \in \mathbb{R}^d$ to the output $\Phi \in \mathbb{R}^q$ whose i 'th coordinate reads as follows

$$\Phi_i(\mathbf{x}|\Theta) := \varphi_i^{(2)} \left(\sum_{j=1}^{p_1} W_{ij}^{(2)} \varphi_j^{(1)} \left(\sum_{k=1}^d W_{jk}^{(1)} x_k + b_j^{(1)} \right) + b_i^{(2)} \right) \quad (18)$$

where $\Theta := (\boldsymbol{\theta}^{(1)}, \boldsymbol{\theta}^{(2)}) := (\mathbf{W}^{(1)}, \mathbf{b}^{(1)}, \mathbf{W}^{(2)}, \mathbf{b}^{(2)}) \in \mathbb{R}^{p_1 \times (d+1) + q \times (p_1+1)}$

- The total number of parameters in the statistical model is

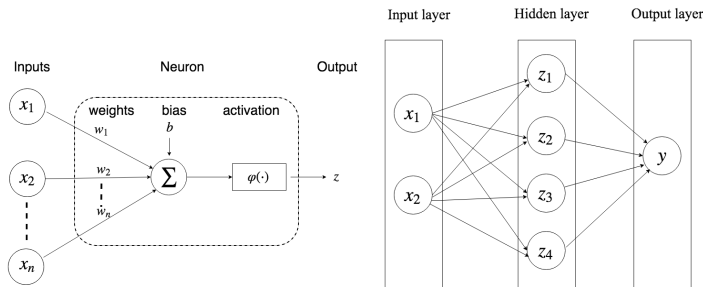
$$p = p_1 \times (d + 1) + q \times (p_1 + 1) \quad (19)$$

(Fully-Connected) Feedforward Artificial Neural Networks

- Denoting the *element-wise* non-linearities by $\varphi^{(n)}(\mathbf{z}^{(n)}) := (\varphi_1^{(n)}(z_1^{(n)}), \dots, \varphi_{p_n}^{(n)}(z_{p_n}^{(n)})) \in \mathbb{R}^{p_n}$, in vector notation this can be written as follows

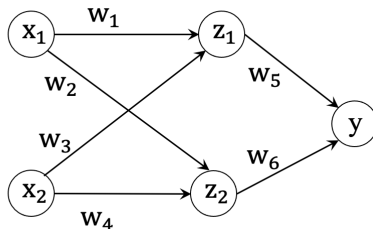
$$\Phi(\mathbf{x}|\Theta) := \varphi^{(2)} \circ \mathbf{z}^{(2)}(\cdot|\theta^{(2)}) \circ \varphi^{(1)} \circ \mathbf{z}^{(1)}(\mathbf{x}|\theta^{(1)}) \quad (20)$$

- FCFF ANNs can be thought of as **directed, acyclic graphs**



Example: XOR revisited

- Unlike with SLP, adding a *hidden layer*, and parametrizing with ANNs solves the problem



- x_1 and x_2 are first mapped to intermediate binary values $z_1 = h_{b_1}(w_1x_1 + w_3x_2)$ and $z_2 = h_{b_2}(w_2x_1 + w_4x_2)$

Example: XOR revisited

- Using $w_1 = -1, w_2 = 1, w_3 = 1, w_4 = -1$, the outputs are shown in columns z_1 and z_2 . Passing this onto the non-linearities h_{b_1}, h_{b_2} with $b_1 = b_2 = 0$ results in activations a_1, a_2 .

x_1	x_2	z_1	z_2	a_1	a_2	y
-1	-1	2	-2	1	-1	1
-1	1	0	0	-1	-1	-1
1	-1	0	0	-1	-1	-1
1	1	-2	2	-1	1	1

- It remains to compute $h_{b_3}(w_5 z_1 + w_6 z_2)$, by setting threshold $b_3 = -1$ and $w_5 = 1, w_6 = 1$ the desired output is obtained in y and the XOR-problem is solved.

Activations: the key feature of non-linearity

Activation Functions

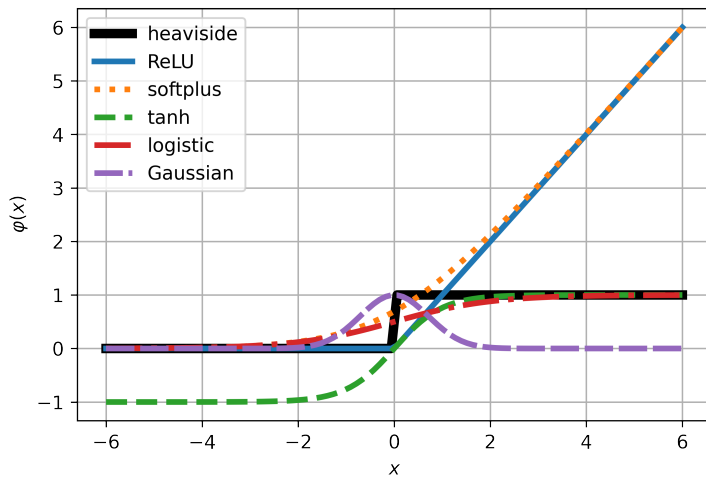
- The key component of *non-linearity* is given by the **activation functions** $\varphi_j^{(n)} : \mathbb{R} \rightarrow \mathbb{R}, j = 1, \dots, p_n$
- For SLP we had $\varphi_1^{(1)}(x) = \text{sgn}(x)$
- Most often, application functions within the same layer are chosen the same $\varphi_1^{(n)}(x) = \dots = \varphi_{p_n}^{(n)}(x)$
- In regression applications the output activations are usually chosen to be the identity function $\varphi^{(2)}(\mathbf{z}^{(2)}) = \mathbf{z}^{(2)}$
- The parametrization inherits **continuity** and **differentiability** properties from the chosen activations
- The non-linear activation functions are *hyperparameters* of the statistical model, appropriate choices vary depending on the problem

Activation Functions

name	$\varphi(x)$	$\varphi'(x)$	range	continuity
heaviside	$\begin{cases} 0, & x \leq 0 \\ 1, & x > 0 \end{cases}$	$\begin{cases} 0, & x \neq 0 \\ \text{undefined}, & x = 0 \end{cases}$	$\{0, 1\}$	C^{-1}
ReLU	$\begin{cases} 0, & x \leq 0 \\ x, & x > 0 \end{cases}$	$\begin{cases} 0, & x < 0 \\ x, & x > 0 \\ \text{undefined}, & x = 0 \end{cases}$	$[0, \infty)$	C^0
Gaussian	e^{-x^2}	$-2xe^{-x^2}$	$(0, 1]$	C^∞
tanh	$\frac{e^x - e^{-x}}{e^x + e^{-x}}$	$1 - \tanh^2(x)$	$(-1, 1)$	C^∞
logistic	$\frac{1}{1 + e^{-x}}$	$\frac{e^{-x}}{1 + e^{-x}}$	$(0, 1)$	C^∞
softplus	$\log(1 + e^x)$	$\frac{1}{1 + e^{-x}}$	$(0, \infty)$	C^∞
...

See wiki

Activation Functions



Layer dependent activations

- For classification problems it is chosen such that $\sum_{k=1}^q \varphi_k^{(2)}(z_k^{(2)}) = 1$
- In order to ensure that output *probabilities* indeed sum to one, activations are chosen such that they are not a collection of scalar-valued mappings of neurons but vector-valued functions depending on the whole layer $\varphi^{(n)} : \mathbb{R}^{p_n} \rightarrow \mathbb{R}^{p_n}$

name	$\varphi_i(\mathbf{x})$	$\partial_j \varphi_i(\mathbf{x})$	range	cont.
softmax	$\frac{e^{x_i}}{\sum_{j=1}^{p_n} e^{x_j}}$	$\varphi_i(\mathbf{x})(\delta_{ij} - \varphi_j(\mathbf{x}))$	$(0, 1)$	C^∞
maxout	$\max_{i=1, \dots, p_n} x_i$	$\begin{cases} 1, & j = \arg \max_i x_i \\ 0, & j \neq \arg \max_i x_i \end{cases}$	$(-\infty, \infty)$	C^0
...

Deep Neural Networks

Fully-Connected Feedforward Deep Neural Networks

- So far: neural network = input + hidden + output layer
- These are called **shallow** feedforward ANNs

$$\Phi(\mathbf{x}|\Theta) := \varphi^{(2)} \circ \mathbf{z}^{(2)}(\cdot|\theta^{(2)}) \circ \varphi^{(1)} \circ \mathbf{z}^{(1)}(\mathbf{x}|\theta^{(1)})$$

- One can extend the model to allow for L -many hidden layers in the composition, with corresponding widths $p_l, l = 1 \dots, L$
- The resulting, **deep neural network** mapping reads as follows

$$\Psi(\mathbf{x}|\Theta) := \varphi^{(L+1)} \circ \mathbf{z}^{(L+1)}(\cdot|\theta^{(L+1)}) \circ \dots \circ \varphi^{(1)} \circ \mathbf{z}^{(1)}(\mathbf{x}|\theta^{(1)}), \quad (21)$$

where $\Theta := (\theta^{(1)}, \dots, \theta^{(L+1)}) \in \mathbb{R}^p$

- The total number of parameters in the statistical model rapidly increases

$$p = (d + 1) \times p_1 + \sum_{l=1}^{L-1} (p_l + 1) \times p_{l+1} + (p_L + 1) \times q \quad (22)$$

Fully-Connected Feedforward Deep Neural Networks

$$\Psi(\mathbf{x}|\Theta) := \varphi^{(L+1)} \circ \mathbf{z}^{(L+1)}(\cdot|\theta^{(L+1)}) \circ \dots \circ \varphi^{(1)} \circ \mathbf{z}^{(1)}(\mathbf{x}|\theta^{(1)}),$$

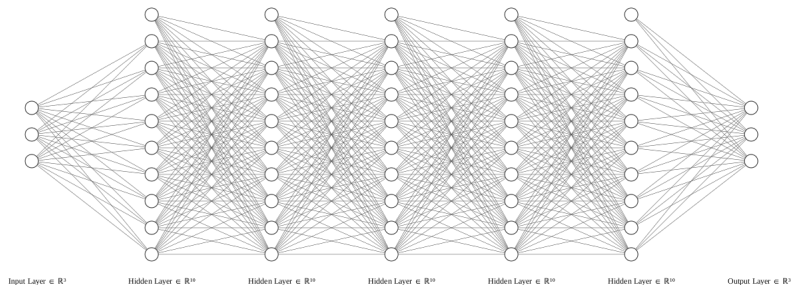


Figure: Illustration DNN, source: NN-SVG

$$p = (d + 1) \times p_1 + \sum_{l=1}^{L-1} (p_l + 1) \times p_{l+1} + (p_L + 1) \times q$$

Teaser: capacity and optimization

What we've seen so far

- We gathered **one potential** parametrization for non-linear problems via (deep) neural networks: $\Phi(\mathbf{x}|\Theta) : \mathbb{R}^d \rightarrow \mathbb{R}^q$
- Similarly to MSE in OLS, one can define a *performance measure* $L(\Theta|\mathbf{x})$ which quantifies the "goodness" of a certain parameter set, under measurements \mathbf{x}
- Such **losses** can be of various forms
 - $L(\Theta|\mathbf{x}) = \mathbb{E}[d(\Phi(\mathbf{x}|\Theta), H(\mathbf{x}))]$ (supervised learning)
 - $L(\Theta|\mathbf{x}) = \mathbb{E}[d(F(\Phi(\mathbf{x}|\Theta)), 0)]$ (unsupervised learning)

where d is a desired *distance function* – such as the ones induce by L^p norms $1 \leq p < \infty$

$$d(f, g) = \|f - g\|_p := \sqrt[p]{\int_{\mathbf{x}} |f(\mathbf{x}) - g(\mathbf{x})|^p d\mu(\mathbf{x})}$$

Questions remaining

- DNNs provide a certain family of statistical models **how do we know what problems are they also appropriate for?**
- In order to find a "good" statistical model we want to minimize the loss function

$$\arg \min_{\theta} L(\theta|\mathbf{x}),$$

- Unlike in case of OLS, the MSE of supervised, neural network regression is no longer *convex* in $\Theta \mapsto L(\Theta|\mathbf{x}) \rightarrow$ **no closed-form minimizer**

Next week

- 1 **Universal Approximation Theorem (UAT):** ANNs are *dense* in a very wide class of target functions \mathcal{H}
- 2 **Stochastic Gradient Descent (SGD):** an empirically successful iterative optimization method